

LECTURE 6

EXECUTION PLAN

- Execution plan includes
 - Access methods available for each relation
 - Algorithms to be used in computing the relational operators represented in the tree
- For materialized evaluation, the result of operation is stored as a temporary relation
- For pipelined evaluation, as the resulting tuples of an operation are produced, they are forwarded directly to the next operation in the query sequence
- Advantage of pipelining – cost savings (No need to write intermediate results to disk)

COST ESTIMATES IN QUERY OPTIMIZATION (1/2)

- More suitable for **compiled queries** where the optimization is done at compile time and the resulting execution strategy code is stored and executed directly at runtime
- For interpreted queries, where entire process occurs at runtime, a full-scale optimization may slow down the response time
- This approach is called **cost-based query optimization**

COST COMPONENTS FOR QUERY EXECUTION (1/3)

1. Access cost to secondary storage

- Cost of searching for reading and writing data blocks that reside on secondary storage (disk)
- Cost of searching records in a file depends on the type of access structures on that file (ordering, hashing, primary or secondary indexes)
- Whether file blocks are allocated contiguously on same disk cylinder or scattered on the disk

2. Storage cost

- Cost of storing any intermediate files that are generated by execution strategy for the query

COST COMPONENTS FOR QUERY EXECUTION (2/3)

3. Computation cost

- Cost of performing in-memory operations on data buffers during query execution
- Includes searching for and sorting records, merging records for a join, and performing computations on field values

4. Memory usage cost

- Cost pertaining to number of memory buffers needed during query execution

5. Communication cost

- Cost of shipping the query and its results from the database site to the site or terminal where the query originated

COST COMPONENTS FOR QUERY EXECUTION (3/3)

- For large database
 - Main emphasis is on minimizing access **cost to secondary storage**
 - Simple cost functions ignore other factors and compare different query execution strategies in terms of the number of block transfers between disk and main memory
- For smaller database
 - Minimizing **computation cost** as most of data in files involved in query can be completely stored in memory
- In distributed database
 - **Communication cost** must be minimized as many sites are involved

CATALOG INFORMATION USED IN COST FUNCTIONS (1/2)

■ Size of the file :

- Number of records (r)
- Record size (R)
- Number of blocks (b)
- Blocking factor (bfr)

■ Access methods and indexes

- No of levels (x) of each multilevel index (primary, secondary, or clustering)
- No of first-level index blocks (b_{l1}) – for some cost functions

CATALOG INFORMATION USED IN COST FUNCTIONS (2/2)

- **Selectivity (sl)** is the fraction of records satisfying an equality condition on the attribute
- **Selection Cardinality (s)** is the average number of records that will satisfy an equality selection condition on that attribute, calculated as **$s = sl * r$**
- **No of distinct values (d)**
- For a key attribute,
 $d = r, sl = 1/r$, and $s = 1$
- For a nonkey attribute, assume that d distinct values are uniformly distributed among the records , we have
 $sl = (1/d)$ hence, $s = (r/d)$

EXAMPLES OF COST FUNCTIONS FOR SELECT (1/8)

- We denote cost for method S_i as C_{si}
- **S1: Linear search**
 - Search all file blocks to retrieve all records satisfying the selection condition
 - Hence, $C_{S1a} = b$
 - For equality condition on a key, only half the file blocks are searched on the average before finding the record

$$C_{S1b} = (b/2) \text{ if record is found}$$

$$C_{S1b} = b \text{ if no record satisfies condition}$$

EXAMPLES OF COST FUNCTIONS FOR SELECT (2/8)

■ S2 : Binary search

- Search accesses approx.

$$C_{s2} = \log_2 b + [(s/bfr)] - 1 \text{ file blocks}$$

- This reduces to $\log_2 b$ if equality condition is on a unique (key) attribute, because $s = 1$ in this case

EXAMPLES OF COST FUNCTIONS FOR SELECT (3/8)

■ S3 : Using primary index (S3a) or hash key (S3b) to retrieve single record

■ For primary index

- Retrieve one more block than the no of index levels

$$C_{S3a} = x + 1$$

■ For hashing

$$C_{S3b} = 1 \text{ for static or linear hashing}$$

$$C_{S3b} = 2 \text{ for extendible hashing}$$

EXAMPLES OF COST FUNCTIONS FOR SELECT (4/8)

- **S4: Using ordering index to retrieve multiple records**
 - If comparison condition is $>, \geq, <, \leq$ on a key field with ordering index, roughly half the file records will satisfy the condition

$$C_{S4} = x + (b/2)$$

EXAMPLES OF COST FUNCTIONS FOR SELECT (5/8)

- **S5: Using clustering index to retrieve multiple records**
 - Given equality condition, s records will satisfy condition, where s is selection cardinality of indexing attribute
 - This means that $\lceil (s/bfr) \rceil$ file blocks will be accessed

$$C_{S5} = x + \lceil (s/bfr) \rceil$$

EXAMPLES OF COST FUNCTIONS FOR SELECT (6/8)

■ S6: Using a secondary index

- On equality comparison, s records will satisfy the condition, where s is selection cardinality of indexing attribute
- Since index is non clustering, each of the records may reside on a different block, hence

$$C_{S6a} = x + s$$

- This reduces to $x+1$ for a key indexing attribute

EXAMPLES OF COST FUNCTIONS FOR SELECT (7/8)

- If comparison condition is $<, >=, <, <=$ and half of the file records are assumed to satisfy the condition, then roughly half the first-level index blocks are accessed plus half the file records via the index

$$C_{s6b} = x + (b_{i1}/2) + (r/2)$$

- $r/2$ factor can be refined if better selectivity estimates are available

EXAMPLES OF COST FUNCTIONS FOR SELECT (8/8)

- **S7: Conjunctive selection**
 - Use either S1 or one of the methods S2 to S6
 - In latter case, we use one condition to retrieve the records and then check in the memory buffer whether each retrieved record satisfies the remaining conditions in the conjunction
- **S8: Conjunctive selection using composite index**
 - Same as S3a, S5, S6a depending on type of index

EXAMPLES OF USING COST FUNCTIONS (1/7)

- EMPLOYEE file has $r_E=10,000$ records stored in $b_E=2000$ disk blocks with blocking factor $bfr_E=5$ records/block and the following access paths :
 1. A clustering index on Salary, with levels $x_{Salary}=3$ and average selection cardinality $s_{Salary}=20$
 2. A secondary index on key attribute Eno with $x_{Eno}=4$ ($s_{Eno}=1$)
 3. A secondary index on the nonkey attribute Dno, with $x_{Dno}=2$ and first level index blocks $b_{l1Dno}=4$. There are $d_{Dno}=125$ distinct values for Dno, so the selection cardinality of Dno is $s_{Dno}=(r_E/d_{Dno})=80$
 4. A secondary index on Sex, with $x_{Sex}=1$. There are $d_{Sex}=2$ values for the sex attribute, so the average selection cardinality is $s_{Sex}=(r_E/d_{Sex})=5000$

EXAMPLES OF USING COST FUNCTIONS (2/7)

- Illustrate the use of cost functions with the examples below

(OP1) : $\sigma_{\text{Eno}='123456789'}$ (EMPLOYEE)

(OP2) : $\sigma_{\text{Dno}>5}$ (EMPLOYEE)

(OP3) : $\sigma_{\text{Dno}=5}$ (EMPLOYEE)

(OP4) : $\sigma_{\text{Dno}=5 \text{ AND Salary}>30000 \text{ AND Sex}='F'}$ (EMPLOYEE)

- Cost of linear search option S1 estimated as

$C_{S1a} = b_E = 2000$ (for nonkey attribute)

$C_{S1b} = (b_E/2) = 1000$ (average cost for selection on key attribute)

EXAMPLES OF USING COST FUNCTIONS (3/7)

(OP1) : $\sigma_{\text{Eno}} = \text{'123456789'}$ (EMPLOYEE)

- Use either S1 ($C_{S1b} = 1000$)

- Method S6

$$\begin{aligned} C_{S6a} &= x_{\text{Eno}} + 1 \\ &= 4 + 1 \\ &= 5 \end{aligned}$$

- Method S6 is chosen

EXAMPLES OF USING COST FUNCTIONS (4/7)

(OP2) : $\sigma_{Dno>5}$ (EMPLOYEE)

- Method S1 ($C_{S1a}=2000$)

- Method S6

$$\begin{aligned}C_{S6b} &= x_{Dno} + (b_{l1Dno}/2) + (r_E/2) \\&= 2 + (4/2) + (10000/2) \\&= 5004\end{aligned}$$

- Method S1 is chosen

EXAMPLES OF USING COST FUNCTIONS (5/7)

(OP3) : $\sigma_{Dno=5}$ (EMPLOYEE)

- Method S1 ($C_{S1a}=2000$)

- Method S6

$$\begin{aligned}C_{S6a} &= x_{Dno} + s_{Dno} \\ &= 2 + 80 \\ &= 82\end{aligned}$$

- Method S6 is chosen

EXAMPLES OF USING COST FUNCTIONS (6/7)

(OP4) : $\sigma_{Dno=5 \text{ AND } Salary>30000 \text{ AND } Sex='F'}$ (EMPLOYEE)

- It has conjunctive selection, estimate the cost of using any one of the three components of selection condition to retrieve records plus brute force approach (linear search)
- Method S1 ($C_{S1a}=2000$)
- For ($Dno=5$), cost estimate is $C_{S6a}=82$
- For ($Salary>30000$), cost estimate is
$$\begin{aligned}C_{S4} &= x_{Salary} + (b_E/2) \\ &= 3 + (2000/2) \\ &= 1003\end{aligned}$$

EXAMPLES OF USING COST FUNCTIONS (7/7)

- For (Sex='F'), cost estimate is

$$\begin{aligned}\text{CS6a} &= x_{\text{Sex}} + s_{\text{Sex}} \\ &= 1 + 5000 \\ &= 5001\end{aligned}$$

- Method S6a for (Dno=5) is chosen
- Hence, condition (Dno=5) is used to retrieve records, and remaining part of conjunctive condition (Salary>30000 AND Sex='F') is checked for each selected record after it is retrieved into memory

EXAMPLES OF COST FUNCTIONS FOR JOIN (1/7)

- Need to have estimate for size (number of tuples) of file that results after JOIN
- Usually kept as a *ratio of the size (number of tuples) of resulting join file to the size of Cartesian product file*, if both are applied to same input files, and is called **join selectivity (js)**
- If number of tuples of a relation R is $|R|$, we have
$$js = |(R \bowtie_c S)| / |(R \times S)| = |(R \bowtie_c S)| / (|R| * |S|)$$
- If no join condition c is there, then $js=1$. The join is same as the CARTESIAN PRODUCT.
- If no tuples from relations satisfy join condition, then $js=0$

EXAMPLES OF COST FUNCTIONS FOR JOIN (2/7)

- Hence, in general $0 \leq j_s \leq 1$

For a join where condition c is equality comparison $R.A=S.B$, we get

Case 1: If A is a key of R , then

$$|(R \bowtie_c S)| \leq |S|, \text{ so } j_s \leq (1/|R|)$$

Case 2: If B is a key of S , then

$$|(R \bowtie_c S)| \leq |R|, \text{ so } j_s \leq (1/|S|)$$

EXAMPLES OF COST FUNCTIONS FOR JOIN (3/7)

- For estimating size of resulting file if size of two input files are known use formula below

$$|(R \bowtie_c S)| = j_s * |R| * |S|$$

- Assume that R has b_R blocks and S has b_S blocks

EXAMPLES OF COST FUNCTIONS FOR JOIN (4/7)

■ J1: Nested-loop join

- Suppose R is outer loop

$$C_{J1} = b_R + (b_R * b_S) + ((js * |R| * |S|) / bfr_{RS})$$

- Last part of formula is cost of writing resulting file to disk

EXAMPLES OF COST FUNCTIONS FOR JOIN (5/7)

■ J2: Single-loop join (using access structure)

- If index exists for join attribute B of S with index levels x_B , we can retrieve each record s in R and then use index to retrieve all matching records t from S that satisfy $t[B]=s[A]$
- Cost depends on type of index
- Secondary index where s_B is selection cardinality for join attribute B of S, we have

$$C_{J2a} = b_R + ((|R| * (x_B + s_B)) + ((js * |R| * |S|) / bfr_{RS}))$$

EXAMPLES OF COST FUNCTIONS FOR JOIN (6/7)

- **Clustering index** where s_B is selection cardinality of B

$$C_{J2b} = b_R + ((|R| * (x_B + (s_B / bfr_B))) + ((js * |R| * |S|) / bfr_{RS}))$$

- **Primary index**

$$C_{J2c} = b_R + ((|R| * (x_B + 1)) + ((js * |R| * |S|) / bfr_{RS}))$$

- **If hash key exists for one of the two attributes, say B of S**

$$C_{J2d} = b_R + (|R| * h) + ((js * |R| * |S|) / bfr_{RS})$$

where $h \geq 1$ is average number of block accesses to retrieve a record, given hash key value

EXAMPLES OF COST FUNCTIONS FOR JOIN (7/7)

■ J3: Sort-merge join

- If sorted on join attributes, cost function is

$$C_{J3a} = b_R + b_S + ((js * |R| * |S|) / bfr_{RS})$$

- If files need to be sorted, then cost of sorting is also added.
- **Sorting cost** is $(2*b) + (2*b*\log_2 b)$ for a file of b blocks

$$C_{J3b} = (2*b_R*(1+\log_2 b_R)) + (2*b_S*(1+\log_2 b_S)) + b_R + b_S + ((js * |R| * |S|) / bfr_{RS})$$

EXAMPLES OF USING COST FUNCTIONS (1/4)

- EMPLOYEE file has $r_E=10,000$ records stored in $b_E=2000$ disk blocks with blocking factor $bfr_E=5$ records/block and the following access paths :
 1. A clustering index on Salary, with levels $x_{Salary}=3$ and average selection cardinality $s_{Salary}=20$
 2. A secondary index on key attribute Eno with $x_{Eno}=4$ ($s_{Eno}=1$)
 3. A secondary index on the nonkey attribute Dno, with $x_{Dno}=2$ and first level index blocks $b_{l1Dno}=4$. There are $d_{Dno}=125$ distinct values for Dno, so the selection cardinality of Dno is $s_{Dno}=(r_E/d_{Dno})=80$
 4. A secondary index on Sex, with $x_{Sex}=1$. There are $d_{Sex}=2$ values for the sex attribute, so the average selection cardinality is $s_{Sex}=(r_E/d_{Sex})=5000$

EXAMPLES OF USING COST FUNCTIONS (2/4)

- We have DEPARTMENT file consisting of $r_D=125$ records stored in $b_D=13$ disk blocks
- There is a primary index on DNumber of DEPARTMENT with $X_{DNumber}=1$ level

EMPLOYEE  Dno=DNumber DEPARTMENT

- The join selectivity for OP6 is $js_{OP6}=(1/|DEPARTMENT|)=1/125$ since DNumber is key of DEPARTMENT
- Assume blocking factor for resulting file is $bfr_{ED}=4$ records per block

EXAMPLES OF USING COST FUNCTIONS (3/4)

1. Using Method J1 with EMPLOYEE as outer loop

$$\begin{aligned}C_{J1} &= b_E + (b_E * b_D) + ((js_{OP6} * r_E * r_D) / bfr_{ED}) \\ &= 2000 + (2000 * 13) + (((1/125) * 10000 * 125) / 4) \\ &= 30500\end{aligned}$$

2. Using method J1 with DEPARTMENT as outer loop

$$\begin{aligned}C_{J1} &= b_D + (b_E * b_D) + ((js_{OP6} * r_E * r_D) / bfr_{ED}) \\ &= 13 + (2000 * 13) + (((1/125) * 10000 * 125) / 4) \\ &= 28513\end{aligned}$$

EXAMPLES OF USING COST FUNCTIONS (4/4)

3. Using method J2 with EMPLOYEE as outer loop

$$\begin{aligned}C_{J2} &= b_E + (r_E * (x_{DNumber} + 1)) + ((js_{OP6} * r_E * r_D) / bfr_{ED}) \\ &= 2000 + (10000 * 2) + (((1/125) * 10000 * 125) / 4) \\ &= 24500\end{aligned}$$

4. Using method J2 with DEPARTMENT as outer loop



$$\begin{aligned}C_{J2} &= b_D + (r_D * (x_{Dno} + s_{Dno})) + ((js_{OP6} * r_E * r_D) / bfr_{ED}) \\ &= 13 + (125 * (2 + 80)) + (((1/125) * 10000 * 125) / 4) \\ &= 12763\end{aligned}$$

- Since case 4 has lowest cost estimate it is chosen

MULTIPLE RELATION QUERIES AND JOIN ORDERING – KINDLY GO THROUGH YOURSELF

- **A query that joins N relations will have $N-1$ join operations**
- Estimating cost of every possible join tree for a query with large number of joins will require huge amount of time by query optimizer
- Query optimizer limits the structure of a (join) query tree to that of left-deep (or right-deep) trees
- **Left-deep tree is a binary tree where right child of each nonleaf node is always a base relation**
- Optimizer will choose the particular left-deep tree with lowest estimated cost

LEFT-DEEP JOIN TREE (1/2) – PLEASE STUDY YOURSELF



LEFT-DEEP JOIN TREE (2/2) – PLEASE STUDY YOURSELF

- With left-deep trees, the **right child is considered to be the inner relation when executing a nested-loop join**
- Advantages of left-deep (or right-deep) tree
 - They are amenable to pipelining
 - Having a base relation as one of the inputs of each join allows the optimizer to utilize any access paths on that relation that may be useful in executing the join