

LECTURE 3

SHADOW PAGING (1/7)

- Shadow paging considers the database to be made up of a number of **fixed-size disk pages** (say n) for recovery purposes
- A **directory** with n entries is constructed, where the i^{th} entry points to the i^{th} database page on disk
- **Directory** is kept in main memory if it is not too large, and all references – reads or writes – to database pages on disk go through it

SHADOW PAGING (2/7)

- When transaction begins executing, the **current directory** – whose entries point to the most recent or current database pages on disk – is copied into a **shadow directory**
- **Shadow directory** is then saved on disk while the **current directory** is used by the transaction
- During transaction execution, **shadow directory** is **never modified**

SHADOW PAGING (3/7)

- When a *write_item* operation is performed, a new copy of the modified database page is created, but old copy is not overwritten
- New page is written elsewhere – on some previously unused disk block
- Current directory is modified to point to the new disk block, whereas shadow directory is not modified and continues to point to the old unmodified disk block

SHADOW PAGING (4/7)

- For pages updated by transaction, **two versions are kept**
 - **Old version is referenced by the shadow directory**
 - **New version is referenced by the current directory**
- To recover from failure during transaction execution, it is sufficient to free the modified database and to discard the current directory
- The state of the database before execution is available through the shadow directory, and that state is recovered by reinstating the shadow directory

SHADOW PAGING (5/7)

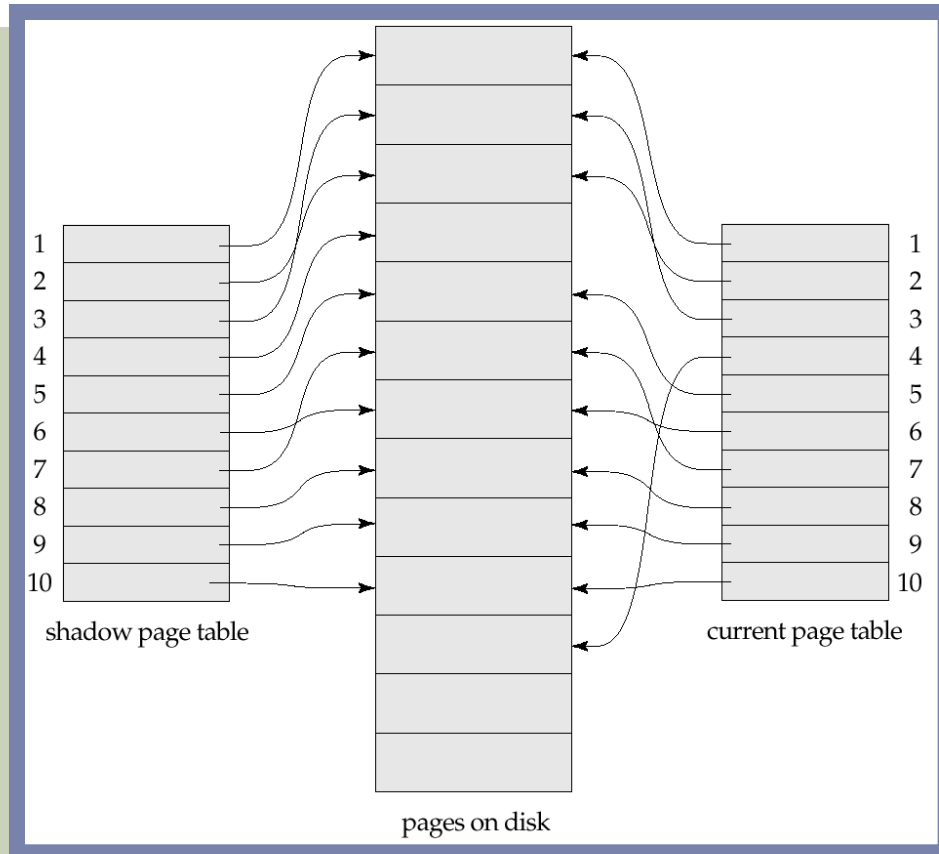
- **Committing a transaction corresponds to discarding the previous shadow directory**
- **This technique can be categorized into NO_UNDO/NO-REDO technique for recovery as neither undoing nor redoing of data items is involved**
- **In multi-user environment with concurrent transactions, logs and checkpoints must be incorporated into shadow paging technique**

SHADOW PAGING (6/7)

■ Disadvantages

- Updated database pages change location on disk. Hence, difficult to keep related database pages close together on disk without complex storage management strategies
- Overhead of writing shadow directories to disk as transactions commit is significant
- Complication on handling garbage collection when a transaction commits. Old pages referenced by shadow directory that have been updated must be released and added to a list of free pages for future use

SHADOW PAGING (7/7)



Example: Shadow and current page tables after write to page 4

RECOVERY WITH CONCURRENT TRANSACTIONS (1/13)

- The system has a single disk buffer and a single log
- All transactions share the buffer blocks
- Allow immediate modification and permit a buffer block to have data items updated by one or more transactions

RECOVERY WITH CONCURRENT TRANSACTIONS (2/13)

■ Interaction with concurrency control

- Recovery scheme depends on concurrency-control scheme that is used
- To roll back a failed transaction, must undo the updates performed by the transaction
- Suppose, if T_0 has to be rolled back, and a data item X that was updated by T_0 has to be restored to its old value
- Restore the value by using undo information in a log record

RECOVERY WITH CONCURRENT TRANSACTIONS (3/13)

- Suppose T_1 has performed another update on X before T_0 is rolled back
- The update performed by T_1 will be lost if T_0 is rolled back
- Hence, if T has updated X , no other transaction may update the same data item until T has committed or been rolled back
- Can ensure this by using strict 2PL, i.e. 2PL with exclusive locks held until the end of the transaction

RECOVERY WITH CONCURRENT TRANSACTIONS (4/13)

■ Transaction rollback

- Rollback a failed transaction T_i by using the log
- System scans the log backward
- For every log record of the form $\langle T_i, X_j, V_1, V_2 \rangle$ found in the log, system restores the data item X_j to its old value V_1
- Scanning of log terminates when the log record $\langle T_i, \text{start} \rangle$ is found

RECOVERY WITH CONCURRENT TRANSACTIONS (5/13)

- Scanning a log backward is important, since a transaction may have updated a data item more than once
- Example
 - $\langle T_1, A, 10, 20 \rangle$
 - $\langle T_1, A, 20, 30 \rangle$
- Scanning log backwards sets A correctly to 10
- If log is scanned forward, then A would be set to 20, which is incorrect

RECOVERY WITH CONCURRENT TRANSACTIONS (6/13)

- If strict 2PL is used for concurrency control, locks held by a transaction T may be released only after T has been rolled back
- Once T has updated a data item, no other transaction could have updated the same data item, because of concurrency control

RECOVERY WITH CONCURRENT TRANSACTIONS (7/13)

■ Checkpoints

- The following transactions must be considered during recovery
 - Transactions that started after the most recent checkpoint
 - The one transaction, if any, that was active at the time of the most recent checkpoint
- In a concurrent transaction processing system, we require that the checkpoint log record be of the form **<checkpoint L>**, where **L** is a list of transactions active at the time of checkpoint

RECOVERY WITH CONCURRENT TRANSACTIONS (8/13)

- Assume that transactions do not perform updates either on buffer blocks or on log while checkpoint is in progress
- **A fuzzy checkpoint is a checkpoint where transactions are allowed to perform updates even while buffer blocks are being written out**
- When system recovers from crash, it constructs two lists
 - **Undo-list** consists of transactions to be undone
 - **Redo-list** consists of transactions to be redone

RECOVERY WITH CONCURRENT TRANSACTIONS (9/13)

- Initially, both lists are empty
- System scans log backward, examining each record, until it finds the first <checkpoint> record
 - For each record found of the form < T_i commit>, it adds T_i to redo-list
 - For each record found of the form < T_i start>, if T_i is not in redo-list, then it adds T_i to undo-list
- When system has examined all the appropriate log records, it checks the list L in checkpoint record

RECOVERY WITH CONCURRENT TRANSACTIONS (10/13)

- For each T_i in L , if T_i is not in redo-list then it adds T_i to undo-list
- Once the two lists have been constructed, recovery proceeds as follows
 1. System rescans log from the most recent record backward, and performs an undo for each log record that belongs to T_i on undo-list. Scan stops when the $\langle T_i, \text{start} \rangle$ records have been found for every T_i in undo-list. Ignores records of redo-list
 2. System locates the most recent $\langle \text{checkpoint } L \rangle$ record on the log. This step may involve scanning log forward, if checkpoint record was passed in step 1

RECOVERY WITH CONCURRENT TRANSACTIONS (11/13)

3. System scans the log forward from the most recent **<checkpoint L>** record, and performs redo for each log record that belongs to T_i that is on the redo-list. It ignores records of undo-list
- Step 1 is important to process log backward to ensure that the resulting state of database is correct
 - After system has undone all transactions on undo-list, it redoes those on the redo-list, hence the log is processed forward

RECOVERY WITH CONCURRENT TRANSACTIONS (12/13)

- When recovery process has completed, transaction processing resumes
- Suppose, data item A has initial value 10, T_i update value of A to 20 and aborted
- Transaction rollback would restore the value of A to 10
- Now, suppose another transaction T_j then updated value of A to 30 and committed, here the system crash

RECOVERY WITH CONCURRENT TRANSACTIONS (13/13)

- State of the log at the time of crash is
 - <T_i, A, 10, 20>
 - <T_j, A, 20, 30>
 - <T_j commit>
- Say, first redo → A set to 30
 second undo → A set to 10 *which is wrong*
- Final value of A should be 30 which we can ensure by performing undo before performing redo